

# HPC Hw3 (Zhe Chen)

## Machine configuration

---

All results are got by running on my office's machine 'blob'. It has a 4 core CPU, Intel Xeon E5-1603, 2.80GHz.

Max memory size is 375GB, max memory bandwidth is 31.4GB/s. It uses Sandy Bridge EP, so operations per cycle is 8. Thus, theoretic FLOPS/s is 4 Cores \* 2.80GHz \* 8 operations/cycle = 89.6 GFLOPs/s

## Makefile

---

All codes are compiled with g++, "O3" flags, openmp. Command is like:

```
g++ -std=c++11 -O3 -march=native -fopenmp target.cpp -o target-omp
```

## P1: Approximating Sine Function with Taylor Series & Vectorization

---

1.First, I improve sin4\_vec to be 12 digit accuracy just by adding higher Taylor terms up to order 11. The performance is shown as below.

### Extra credit

2.To evaluate sine function outside of  $[-\pi/4, \pi/4]$  efficiently, I develop scheme as following.

2.a). Move any 'x' to  $[-\pi/4, 7\pi/4]$  since it's  $2\pi$  circulant by  $x = x - 2\pi \text{floor}(x + \pi/4)$ .

2.b). By  $\sin(x) = -\sin(x - \pi)$ , move  $x \in [3\pi/4, 7\pi/4]$  to  $[-\pi/4, 3\pi/4]$ .

2.c). We can already compute  $\sin(x)$  in  $[-\pi/4, \pi/4]$ . To evaluate it in  $[\pi/4, 3\pi/4]$ , we use formula  $\sin(x) = \cos(x - \pi/2)$ . Thus, the problem is converted to compute  $\cos(x)$ ,  $x \in [-\pi/4, \pi/4]$ . We could just use cosine function's taylor expansion up to  $x^{12}$  to get 12 digits accuracy.

The result is shown in below as "(extended range)".

```
Reference time: 22.9906
Reference time (extended range): 32.7627
Taylor time:    4.4044      Error: 6.928125e-12
Taylor time (extended range): 19.5440      Error: 6.928014e-12
Intrin time:    1.3712      Error: 2.454130e-03
Vector time:    1.5247      Error: 6.928125e-12
```

## P2: Parrallel Scan in OpenMP

---

To parallelize scan operation, I divide  $n$  into  $n_{\text{threads}}$  batches with each batch's length= $n/n_{\text{threads}}$ . Generally, it's better to have  $n_{\text{threads}}$  less or equal to max core number, which is 4 in my machine.

In each batch, we calculate prefix-sum locally and record the last term  $batch[i]$ . Then we scan  $batch[i]$  to get  $prefix\_sum\_batch[i]$ , which is the value that should be added into batch  $i+1$ .

Importantly, we should add  $prefix\_sum\_batch[i]$  to each batch parallelingly too. Or we lose the point of paralleling scan operator. It's also a complete scan to do this operation!

The result is shown below and we can see good scaling.

Number of Threads	1	2	3	4
Parallel Scan / Sequential Scan	1.4336	0.8748	0.6125	0.5047