

HPC Hw5 (Zhe Chen)

Machine configuration

Results are obtained by running on CIMS's HPC cluster Prince.

Modules environment is

```
3) cuda/8.0.44 7) gcc/6.3.0 8) intel/17.0.1 9) openmpi/intel/1.8.8
```

I am only able to request only 16 nodes and 1GB memory since it's too crowded on Prince.

```
srun --nodes=16 --ntasks-per-node=1 -t2:00:00 --mem=10000 --pty /bin/bash
```

P0: Final Project Update

Date	Members	Tasks
04/29 - 05/03	(Zhe & Guanchun)	Literature search, discuss algorithms and write pseudocode and think about implementation with CUDA
05/03 - 05/10	(Zhe & Guanchun)	Complete psedo-code and serial coding and is now trying switching it to parallel implementation

P1: MPI-Parallel 2D Jacobian smoother

I implemented a MPI version of Jacobian methon in Jacobi-omp.cpp. Number of processors should be $p=4^j$ and number of pts is $N=2^j NI$.

One should notice that it should be paid attention to communication order of OMP_Send and OMP_Recv to avoid deadlock. Moreover, it's better to gather a array for communication rather than communicate point-by-point between nodes.

Here we fix number of iteration to be 1000 for the convenience of comparison.

Here's makefile for compiling (-O3 flag) and one should run it as,

```
mpirun -n p ./jacobi-mpi N MaxIters
```

1. Weak scaling: First, we fix $N=100$ and increase j .

Results are showed as below. Increase of consuming time shows the cost of communication between nodes.

j	0	1	2
p	1	4	16
N	100	200	400
t/s	0.017996	0.029307	0.087737

2. Strong scaling: Then, We fix number of points to be $N=400$ and increase nodes to see its scaling with processors.

p	1	4	16
t/s	1.410616	0.152822	0.088982

We got a nealy half scaling as p increase from 4 to 16. However, consuming time for $p=1$ is much bigger than $4 \cdot t(p=4)$, which shows that it's memery bounded as $p=1$.

P2: Parallel sample sort.

I implemented sample sort using parallel algorithm in `ssort.cpp`. Suppose we have P processors, algorithm is as following.

We initialize N uniformly distributed random integers in each processors. A local naive sort is done at first. Then we gathers $p-1$ splitters from each nodes and select $p-1$ splitters uniformly from these $p(p-1)$ number and broadcast them to each nodes. Each node using the splitters to get p bins. One should notice that we only need `lower_bound` function since it's already well-sorted locally. Finally we communicate each bins to its accordingly nodes and write output file locally with filename `"output{p}.txt"`. The results are as following. The scaling of N is good but bad for number of processors p . This is because N is number for each processor. Thus, there're NP number to be sorted in total and cost is paid to communicate between processors.

N;p	4	8	16
10 ⁴	0.022499	0.049597	0.191107
10 ⁵	0.025785	0.060652	0.231230
10 ⁶	0.129566	0.160658	0.319016